

Logistic Function

Example 6-1: The logistic function in Python for one independent variable

```
In [111... import math


def predict_probability(x, b0, b1):
    p = 1 / (1 + math.exp(-(b0 + b1 * x)))
    return p
```

Example 6-2: Using Sympy to plot a logistic function

```
In [112... import math
from sympy import *

x, b0, b1 = symbols('x b0 b1')
p = 1 / (1 + exp(-(b0 + b1 * x)))
p = p.subs(b0, -2.823)
p = p.subs(b1, 0.620)

plot(p); # the plot is from sympy, it can handle both 2D and 3D (plot3d)
```



Using Scipy

Example 6-3. Using a plain logistic regression in Scipy

```
In [112... import pandas as pd
from sklearn.linear_model import LogisticRegression

#Load the data
df = pd.read_csv('https://bit.ly/33ebs2R', delimiter=',')

X = df.values[:, :-1]
Y = df.values[:, -1]

model = LogisticRegression(penalty=None)
model.fit(X,Y)
```

```
Out[112... LogisticRegression
LogisticRegression(penalty=None)
```

```
In [112... b1 = model.coef_.flatten()[0]
```

```
In [113... b0 = model.intercept_[0]
```

Using Maximum Likelihood and Gradient Descent

Example 6-4: Calculating the joint likelihood of observing all the points for a given logistic regression

```
In [113... # Using if functions
import math
import pandas as pd

patient_data = pd.read_csv('https://bit.ly/33ebs2R', delimiter=',').itertupl

def logistic_function(x):
    p = 1 / (1 + math.exp(-(b0 + b1 * x)))
    return p

joint_likelihood = 1

for p in patient_data:
    if p.y == 1:
        joint_likelihood *= logistic_function(p.x)
    elif p.y == 0:
        joint_likelihood *= (1 - logistic_function(p.x))

print(joint_likelihood)
```

4.79111802216874e-05

```
In [113... # Using multiply
import math
import pandas as pd

patient_data = pd.read_csv('https://bit.ly/33ebs2R', delimiter=',').itertupl

def logistic_function(x):
    p = 1 / (1 + math.exp(-(b0 + b1 * x)))
    return p

joint_likelihood = 1

for p in patient_data:
    joint_likelihood *= (logistic_function(p.x) ** p.y * (1 - logistic_funct
#     if p.y == 1:
#         joint_likelihood *= logistic_function(p.x)
#     elif p.y == 0:
#         joint_likelihood *= (1 - logistic_function(p.x))
```

```
print(joint_likelihood)
```

4.79111802216874e-05

```
In [114... # Using log functions - avoiding the floating point underflow
import math
import pandas as pd

patient_data = pd.read_csv('https://bit.ly/33ebs2R', delimiter=',').itertupl

def logistic_function(x):
    p = 1 / (1 + math.exp(-(b0 + b1 * x)))
    return p

joint_likelihood = 0

for p in patient_data:
    if p.y == 1:
        joint_likelihood += math.log(logistic_function(p.x))
    elif p.y == 0:
        joint_likelihood += math.log(1 - logistic_function(p.x))

print(math.exp(joint_likelihood))
```

4.791118022168739e-05

```
In [114... # Calculate the logarithmic addition
import math
import pandas as pd

patient_data = pd.read_csv('https://bit.ly/33ebs2R', delimiter=',').itertupl

def logistic_function(x):
    p = 1 / (1 + math.exp(-(b0 + b1 * x)))
    return p

joint_likelihood = 0.0

for p in patient_data:
    joint_likelihood += math.log(logistic_function(p.x) ** p.y * \
                                (1.0 - logistic_function(p.x)) ** (1.0 - p.y))

joint_likelihood = math.exp(joint_likelihood)
print(joint_likelihood)
```

4.791118022168739e-05

Example 6-8: Using gradient descent on logistic regression

```
In [114... from sympy import *
```

```

import pandas as pd

points = list(pd.read_csv('https://tinyurl.com/y2coco07').itertuples())

b1, b0, i, n = symbols('b1 b0 i n')
x, y = symbols('x y', cls=Function)

joint_likelihood = Sum(log((1.0 / (1.0 + exp(-(b0 + b1 * x(i))))) ** y(i) \
    * (1.0 - (1.0 / (1.0 + exp(-(b0 + b1 * x(i))))) ** (1 - y(i))), (i,

d_b1 = diff(joint_likelihood, b1) \
    .subs(n, len(points)-1).doit() \
    .replace(x, lambda i: points[i].x) \
    .replace(y, lambda i: points[i].y)

d_b0 = diff(joint_likelihood, b0) \
    .subs(n, len(points)-1).doit() \
    .replace(x, lambda i: points[i].x) \
    .replace(y, lambda i: points[i].y)

d_b1 = lambdify([b1, b0], d_b1)
d_b0 = lambdify([b1, b0], d_b0)

b1 = 0.01
b0 = 0.01
L = 0.01

for j in range(10_000):
    b1 += d_b1(b1, b0) * L
    b0 += d_b0(b1, b0) * L

print(b1, b0)

```

0.6926693075370819 -3.1757515504098244

Multivariable Logistic Regression

Example 6-9: Doing a multivariable logistic regression on employee data

```

In [115... import pandas as pd
from sklearn.linear_model import LogisticRegression

employee_data = pd.read_csv("https://tinyurl.com/y6r7qjrp")

inputs = employee_data.iloc[:, :-1]
output = employee_data.iloc[:, -1]

fit = LogisticRegression(penalty=None).fit(inputs, output)

```

```
print(employee_data.columns)
print('coefficients:{0}'.format(fit.coef_.flatten()))
```

```
Index(['SEX', 'AGE', 'PROMOTIONS', 'YEARS_EMPLOYED', 'DID_QUIT'], dtype='object')
```

```
coefficients: [ 0.03213405  0.03682453 -2.50410028  0.9742266 ]
```

```
In [116... # Interact and test with new employee data
def predict_employee_will_stay(sex, age, promotions, years_employed):
    prediction = fit.predict([[sex, age, promotions, years_employed]]) # due
    probabilities = fit.predict_proba([[sex, age, promotions, years_employed]])
    if prediction == [[1]]:
        return 'will leave: {0}'.format(probabilities)
    else:
        return 'will stay: {0}'.format(probabilities)
```

```
In [116... # Test a prediction
n = input("Predict employee will stay or leave {sex}, {age}, {promotions}, {years_employed} ")
(sex, age, promotions, years_employed) = n.split(',')
print(predict_employee_will_stay(int(sex), int(age), int(promotions), int(years_employed)))
```

```
will leave: [[0.28570264 0.71429736]]
```

```
/opt/anaconda3/lib/python3.11/site-packages/sklearn/base.py:439: UserWarning: X does not have valid feature names, but LogisticRegression was fitted with feature names
  warnings.warn(
/opt/anaconda3/lib/python3.11/site-packages/sklearn/base.py:439: UserWarning: X does not have valid feature names, but LogisticRegression was fitted with feature names
  warnings.warn(
```

R-Squared

1) Using statsmodel to get the R_Squared

```
In [117... import statsmodels.api as sm
#Load the data
df = pd.read_csv('https://bit.ly/33ebs2R', delimiter=',')
```

```
In [120... x = df.values[:, :-1]
y = df.values[:, -1]
x = sm.add_constant(x)

# fit the logistic regression model
model = sm.Logit(y, x)
result = model.fit()
```

```
# Get the Logistic Regression Results
result.summary()
```

Optimization terminated successfully.
 Current function value: 0.473627
 Iterations 6

Out [120...

Logit Regression Results

Dep. Variable:	y	No. Observations:	21
Model:	Logit	Df Residuals:	19
Method:	MLE	Df Model:	1
Date:	Sun, 26 May 2024	Pseudo R-squ.:	0.3065
Time:	21:05:47	Log-Likelihood:	-9.9462
converged:	True	LL-Null:	-14.341
Covariance Type:	nonrobust	LLR p-value:	0.003029

	coef	std err	z	P> z	[0.025	0.975]
const	-3.1758	1.553	-2.045	0.041	-6.220	-0.131
x1	0.6927	0.297	2.331	0.020	0.110	1.275

2) Plain Python to calculate the R_Squared

```
In [121... import pandas as pd
import math
#Load the data
points = list(pd.read_csv('https://bit.ly/33ebs2R', delimiter=',').itertuple
```

Calculate the log likelihood fit

```
In [122... # - Method 1
## define the likelihood_fit
def likelihood(x):
    p = 1 / (1 + math.exp(-(b0 + b1 * x)))
    return p

log_likelihood_fit = 0

for p in points:
    if p.y == 1:
        log_likelihood_fit += log(likelihood(p.x))
    elif p.y == 0:
        log_likelihood_fit += log(1 - likelihood(p.x))
```

```
log_likelihood_fit
```

Out [122...] $\displaystyle -9.94616167318397$

```
In [122...] # - Method 2

## define the likelihood_fit
def likelihood(x):
    p = 1 / (1 + math.exp(-(b0 + b1 * x)))
    return p

log_likelihood_fit = 0

log_likelihood_fit = sum(log(likelihood(p.x) ** p.y) + log((1 - likelihood(p.x) ** (1 - p.y))))

log_likelihood_fit
```

Out [122...] $\displaystyle -9.94616167318397$

```
In [126...] # - Method 3

## define the likelihood_fit
def likelihood(x):
    p = 1 / (1 + math.exp(-(b0 + b1 * x)))
    return p

log_likelihood_fit = 0

log_likelihood_fit = sum(log(likelihood(p.x)) * p.y + log(1 - likelihood(p.x) ** (1 - p.y)))

log_likelihood_fit
```

Out [126...] -9.946161673231583

Calculate the log likelihood

```
In [122...] # - Method 1

likelihood = sum(p.y for p in points) / len(points)

log_likelihood = 0

for p in points:
    if p.y == 1:
        log_likelihood += log(likelihood)
    elif p.y == 0:
        log_likelihood += log(1 - likelihood)
```



```
log_likelihood
```

Out [122...] $\displaystyle -14.3410701987099$

```
In [122...] # - Method 2
likelihood = sum(p.y for p in points) / len(points)

log_likelihood = 0

for p in points:
    log_likelihood += log(likelihood) * p.y + log(1 - likelihood) * (1 - p.y)
log_likelihood
```

Out [122...] $\displaystyle -14.3410701987099$

```
In [122...] # - Method 3
likelihood = sum(p.y for p in points) / len(points)

log_likelihood = 0

log_likelihood = sum(log(likelihood) * p.y + log(1 - likelihood) * (1 - p.y))
log_likelihood
```

Out [122...] $\displaystyle -14.3410701987099$

```
In [126...] # - Method 4
likelihood = sum(p.y for p in points) / len(points)

log_likelihood = 0

log_likelihood = sum(log(likelihood ** p.y) + log((1 - likelihood) ** (1 - p.y)))
log_likelihood
```

Out [126...] -14.341070198709906

Calculate the R_Squared

```
In [123...] r_squared = (log_likelihood - log_likelihood_fit) / log_likelihood
print('{0:.4f}'.format(r_squared))
```

0.3065

```
In [123...] # compare with the statsmodel result
print('{0:.4f}'.format(result.prsquared))
```

0.3065

P_Value

```
In [126... from scipy.stats import chi2
```

```
In [126... chi2_input = 2 * (log_likelihood_fit - log_likelihood)
```

```
In [126... p_value
```

```
Out[126... <scipy.stats._distn_infrastructure.rv_continuous_frozen at 0x318248410>
```

```
In [126... log_likelihood_fit
```

```
Out[126... -9.946161673231583
```

```
In [126... log_likelihood
```

```
Out[126... -14.341070198709906
```

```
In [126... chi2_input
```

```
Out[126... 8.789817050956646
```


```
In [127... chi2.pdf(chi2_input,1)
```

```
Out[127... 0.0016604875618753787
```

Check Chi-Square distribution is the square of each value in standard normal distribution


```
In [130... a = np.random.standard_normal(10000)
```

```
In [130... plt.hist(a,bins=50);
```




```
In [130... b = a ** 2
```

```
In [130... plt.hist(b, bins=100);
```




```
In [130... import numpy as np
import matplotlib.pyplot as plt
from scipy.stats import chi2
```


```
In [131... df = 10
x = np.linspace(0,20,1000)
pdf = chi2.pdf(x, df)
plt.plot(pdf);
```



```
In [131... df = 5
x = np.linspace(0,20,1000)
pdf = chi2.pdf(x, df)
plt.plot(pdf);
```



```
In [131... df = 2  
x = np.linspace(0,20,1000)  
pdf = chi2.pdf(x, df)  
plt.plot(pdf);
```



Train/Test Splits

```
In [131... import pandas as pd
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import KFold, cross_val_score
```

```
In [132... df = pd.read_csv('https://tinyurl.com/y6r7qjrp', sep=',')

x = df.values[:, :-1]
y = df.values[:, -1]

kfold = KFold(n_splits=3, random_state=7, shuffle=True)
model = LogisticRegression(penalty=None)
result = cross_val_score(model, x, y, cv=kfold)
result.mean() # Accuracy
```

```
Out[132... 0.6111111111111112
```

Confusion Matrices

```
In [138... import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.datasets import make_classification
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import confusion_matrix, ConfusionMatrixDisplay, classif

# Generate synthetic binary classification data
X, y = make_classification(n_samples=1000, n_features=20, random_state=42)


# Split data into training and test sets, using stratify to solve the imbalance
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)

# Fit a logistic regression model
model = LogisticRegression()
model.fit(X_train, y_train)

# Predict on the test set
y_pred = model.predict(X_test)
```

```
In [138... # Compute the confusion matrix
cm = confusion_matrix(y_test, y_pred)

# Display the confusion matrix
disp = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=model.classes_)
disp.plot(cmap='Blues')
plt.title('Confusion Matrix')
plt.show()
```

```
In [138... model.classes_
```

```
Out[138... array([0, 1])
```

```
In [138... cm
```

```
Out[138... array([[129, 21],  
        [ 20, 130]])
```

```
In [138... cm.sum() # Total variables
```

```
Out[138... 300
```

```
In [138... cm.sum(axis=1) # Get the supports
```

```
Out[138... array([150, 150])
```

```
In [138... # Print classification report  
print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.87	0.86	0.86	150
1	0.86	0.87	0.86	150
accuracy			0.86	300
macro avg	0.86	0.86	0.86	300
weighted avg	0.86	0.86	0.86	300

```
In [138... TP = cm[1,1]
TN = cm[0,0]
FP = cm[0,1]
FN = cm[1,0]
```

```
In [139... # This is also be the precision for 0
Negative_predict_value = TN / (TN + FN)
print('{0:.2f}'.format(Negative_predict_value))
```

0.87

```
In [139... Precision = TP / (TP + FP)
print('{0:.2f}'.format(Precision))
```

0.86

```
In [139... Recall = TP / (TP + FN) # Also called Sensitivity
print('{0:.2f}'.format(Sensitivity))
```

0.83

```
In [139... # This is also be the recall for 0
Specificity = TN / (TN + FP)
print('{0:.2f}'.format(Specificity))
```

0.86

```
In [139... f1_score = 2 * Precision * Recall / (Precision + Recall)
print('{0:.2f}'.format(f1_score))
```

0.86

```
In [139... Accuracy = (TP + TN) / (TP + TN + FP + FN)
print('{0:.2f}'.format(Accuracy))
```

0.86

ROC (Receiver Operator Characters) & AUC (Area Under Curve)

```

In [139... import numpy as np
import pandas as pd
import matplotlib.pyplot as plt

from sklearn.datasets import make_classification
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import train_test_split
from sklearn.metrics import roc_curve, roc_auc_score, RocCurveDisplay

In [140... x, y = make_classification(n_samples=1000, n_features=20, n_classes=2, random_state=42)
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=1/3, random_state=42)
model = LogisticRegression().fit(x_train, y_train)

In [141... y_pred = model.predict(x_test)
y_prob = model.predict_proba(x_test) # returns an array where each column represents the probability of each class
y_prob = y_prob[:,1] #only retrieve the probability for 1 (since the model classifies between 0 and 1)

In [141... fpr, tpr, thresholds = roc_curve(y_test, y_prob)
# fpr = 1 - precision
# tpr = recall
# thresholds: the predicted probability is converted to a binary classification
# examine how the True Positive Rate (TPR) and false Positive Rate (FPR) change with the threshold

In [141... roc_auc = roc_auc_score(y_test, y_prob)

In [142... plt.figure(figsize=(10, 6))
plt.plot(fpr, tpr, color='blue', lw=2, label=f'ROC curve (area = {roc_auc:.2f})')
plt.plot([0,1],[0,1], color='gray', lw=2, linestyle='--', label='Random Guess')
plt.xlim([0,1])
plt.ylim([0,1.05])
plt.xlabel('FPR')
plt.ylabel('TPR')
plt.title('Receiver Operating Characteristic (ROC) Curve')
plt.legend(loc='lower right')
plt.grid()
plt.show();

```

